
COMP90025 ASSIGNMENT 02A WRITTEN REPORT

Xulin Yang
904904
The University of Melbourne
xuliny@student.unimelb.edu.au

September 12, 2020



Figure 1: Matrix cell dependency

As we have the cell calculation dependencies shown in figure 1.

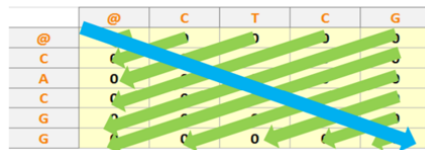


Figure 2: Anti-diagonal matrix traverse

First intuitive optimization is to traverse the matrix in an anti-diagonal manner in figure 2 and parallel each diagonal rather than calculate cells in sequential row by row manner. However, this achieves no speedup because there can be a lot of cache missing and false sharing due to the way we traverse the matrix.

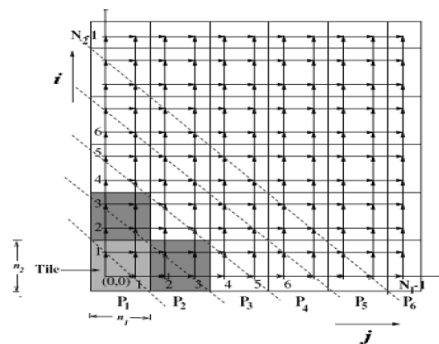


Figure 3: Tiled matrix & Diagonal tile parallelism

Sathe and Shrimankar (2011) describes a tiling parallel algorithm. Rather than parallel each cell diagonally, we parallel **tiles** diagonally. Each tile has $\frac{\text{matrix width}}{p} * \frac{\text{matrix height}}{p}$ cells where $p = \text{\#processors}$ as shown in figure 3. Then

we apply the anti-diagonal parallelism technique on tiles and calculate cells sequentially in each tile to reduce cache missing. Implementation can be found in figure 5. On spartan, I found that 22 threads performs best with this algorithm.

Other speedup techniques:

1. add "#pragma omp parallel for" for the two matrix initialization for loops without implicit barriers in figure 4.
2. delete unnecessary "memset (dp[0], 0, size);" in given code.
3. delete "delete[] dp[0]; delete[] dp;" to speedup.
4. no need to do last branching condition as it needs to execute that step in figure 6.
5. Rather than choose $p = \#processors$, choose $p = \frac{4}{3} * \#processors$ to maximize threads usage.

```
#pragma omp parallel
{
    #pragma omp for nowait
    for (i = 0; i <= m; i++) {
        dp[i][0] = i * pgap;
    }
    #pragma omp for
    for (i = 1; i <= n; i++) {
        dp[0][i] = i * pgap;
    }
}
```

Figure 4: Matrix initialization parallelism

```
167 // calculating the minimum penalty
168
169 // Tile parallel
170 int n_parallel = n_threads + 7;
171 // calculate tile size
172 int tile_width = (int) ceil((1.0*m) / n_parallel), tile_length = (int) ceil((1.0*n) / n_parallel);
173 int num_tile_in_width = (int) ceil((1.0*m) / tile_width);
174 int num_tile_in_length = (int) ceil((1.0*n) / tile_length);;
175
176 // There will be tile_width + num_tile_in_length-1 lines in the output
177 for (int line = 1; line <= (num_tile_in_width + num_tile_in_length - 1); line++) {
178     /* Get column index of the first element in this line of output.
179     The index is 0 for first tile_width lines and line - tile_width for remaining
180     lines */
181     int start_col = max(0, line - num_tile_in_width);
182
183     /* Get count of elements in this line. The count of elements is
184     equal to minimum of line number, num_tile_in_length-start_col and num_tile_in_width */
185     int count = min(line, min((num_tile_in_length - start_col), num_tile_in_width));
186
187     // parallel each tile
188     #pragma omp parallel for
189     for (int z = 0; z < count; z++) {
190         int tile_i_start = (min(num_tile_in_width, line)-z-1)*tile_width +1,
191             tile_j_start = (start_col+z)*tile_length +1;
192
193         // sequential calculate cells in tile
194         for (int i = tile_i_start; i < min(tile_i_start + tile_width, row); i++) {
195             for (int j = tile_j_start; j < min(tile_j_start + tile_length, col); j++) {
196
197                 if (x[i - 1] == y[j - 1]) {
198                     dp[i][j] = dp[i - 1][j - 1];
199                 } else {
200                     dp[i][j] = min3(dp[i - 1][j - 1] + pxy ,
201                                     dp[i - 1][j] + pgap ,
202                                     dp[i][j - 1] + pgap);
203                 }
204             }
205         }
206     }
207 }
208
```

Figure 5: Anti-diagonal tile parallelism & sequential traverse in tile

References

S. R. Sathe and D. D. Shrimankar. Parallelization of DNA sequence alignment using OpenMP. *ACM International Conference Proceeding Series*, pages 200–203, 2011. doi: 10.1145/1947940.1947983.

```

while (!(i == 0 || j == 0)) {
    if (x[i - 1] == y[j - 1]) {
        xans[xpos--] = (int) x[i - 1];
        yans[ypos--] = (int) y[j - 1];
        i--;
        j--;
    } else if (dp[i - 1][j - 1] + pxy == dp[i][j]) {
        xans[xpos--] = (int) x[i - 1];
        yans[ypos--] = (int) y[j - 1];
        i--;
        j--;
    } else if (dp[i - 1][j] + pgap == dp[i][j]) {
        xans[xpos--] = (int) x[i - 1];
        yans[ypos--] = (int) '_';
        i--;
    } // } else if (dp[i][j - 1] + pgap == dp[i][j]) {
    } else {
        xans[xpos--] = (int) '_';
        yans[ypos--] = (int) y[j - 1];
        j--;
    }
}

```

Figure 6: Backtracking optimization